1. CUDA □□ □□□□□ □□

1.1. □□

SM_2D□ □□□□□ □□□□ □□□□ □□ 6□□ CUDA □□ □□□□□□□ □□□□□. □□□□□□□ □□□□ □□□□□ □□ □□□□ □□□□□, □□□□□ □□□□ □□ □□□□ □□ □ □□□□ □□ □□□□ □□ □□□□ □□□□□.

1.1.1. □□ □□□□

□□□□□□ □□□□ □ □□□□□ 6□ □□□□ □□□□□□:

□□	□□
K1: ActiveMask	□□ □□□□ □□□□ □ □□
K2: CoarseTransport	□□□□□ □□(E > 10 MeV)
K3: FineTransport	□□□□□ □□(E <= 10 MeV)
K4: BucketTransfer	□ □ □□ □□
K5: ConservationAudit	□□ □□ □□
K6: SwapBuffers	□□/□□ □□□□ □□

Table 1: CUDA □□ □□□□□□

1.1.2. □□□□□ □□

□□	□□	□□
□□	□□: □□ □□□□ (t = 0)	
K1	ActiveMask □□ □□□□	□□: □□ □□□□ □□□□ □ □□ • □ □□ □□ □□□□ □□ • E < 10 MeV□□ "□□"□□ □□ • □□: ActiveMask □□
K2	CoarseTransport □□ □□	□□: □□□□ □ (E > 10 MeV) • □□ □□□□ □□□□ □□ • □□ □□, □□□□ □ □ • □ □ □□ □□ • □□ (3-5□ □□ □□)
K3	FineTransport □□ □□	□□: □□ □ (E <= 10 MeV) • □□□□ □□ □□□□ • MCS □□ □□□□ • □ □□ □□ □□ • □□□□ (1 percent □□ □□)
K4	BucketTransfer □□ □□	□□: □□ □□ □□ □□ • 4□ □□ □□ (±x, ±z) • □ □□□□ □□ □□□□ □□ □□ • □□□□ □□ □□ □□□□ □□
K5	ConservationAudit □□ □□	□□: □□□□ □□□□□ □□□□□ □□□□□ □□ • □ □□ □□□□ □□ • □ □□ □□□□ □□ • □□: □□ = □□ + □□ • □□ □□
K6	SwapBuffers □□ □□	□□: □□ □□ □□ □□ • □□/□□ □□□□ □□ (CPU □) • □□ □□□□ □□ □□□□ □ • 2.2 GB □□□□ □□ □□
□□	□□: t = Δt□ □□	□□ □□ □□ □□

Table 2: CUDA □□□□□□ □□ □□

1.2. K1: ActiveMask □□

1.2.1. □□

`src/cuda/kernels/k1_activemask.cu`

1.2.2. □□□□ □□

K1□ □□□□□□ □□□□□ □□ □□ □□□□□ □□□□□ □□(E <= 10 MeV)□ □□ □□ □□□□□□. □□□□(0□ 1□ □□)□ □□□□□ □□ □□□□ □□□□ □□ □□□□□□.

1.2.3. □ □□□□

□□□□ □□(E > 10 MeV)□ □□□ □□□□ □□ □□ □□□ □ □□□ □□□ □□□□. □□□□ □□(E <= 10 MeV)□ □□□ □□ □□□ □□□ □□□ □□ □ □□ □ □□ □□□ □□□□□. K1□ □□ □□ □□□□ □□□□ □□ □□□□ □□□□□ □□□□□.

1.2.4. □□ □□

□□	□□	□□
□□	□□ □□ □□ □□□	
□□	□ □□□□ □□□ □ □□: • □□ □□ □□ □□□ □□ • □□ □□□ □□ • 10 MeV □□□□ □□	□□ □□
□□	• □□ E > 10 MeV: ActiveMask = 0 • □□ E <= 10 MeV: ActiveMask = 1 • □□□□ □□ (1e-12 □□)	□□□ □□
□□	• ActiveMask □□ (0 □□ 1) • ActiveList (□□□ □□ □ □□)	K3□ □□

Table 3: K1 □□ □□□ □□ □□

1.2.5. □□□ □□

□□□	□
□□□ □□	`(Nx * Nz + 255) / 256`
□□ □□	256 □□□
□□ □□□	1 □□□ 1 □ □□
□□□ □□	block_ids_in, values_in□ □□ □□

Table 4: K1 □□□ □□

1.2.6. □□□□

```
__global__ void k1_activemask(
    // □□ □□ □□
    const uint32_t* __restrict__ block_ids_in,
    const float* __restrict__ values_in,

    // □□□ □□□□
    const int Nx, const int Nz,

    // □□□
    const float b_E_trigger,      // □□□ □□□ (□□□: 10 MeV)
    const float weight_active_min, // □□ □□□ (□□□: 1e-12)

    // □□
    uint8_t* __restrict__ ActiveMask
);
```

1.2.7. □□□□

```
__global__ void k1_activemask(...) {
    // □□ 1: □ □□□□ □□□□ □ □□
```

```
    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    if (cell >= Nx * Nz) return;  // □□ □□

    // □□ 2: □□ □□ □□□
    float total_weight = 0.0f;
    bool has_low_energy = false;

    // □□ 3: □ □□ □□ □□ □□ □□
    for (int slot = 0; slot < Kb; ++slot) {
        uint32_t bid = block_ids_in[cell * Kb + slot];
        if (bid == EMPTY_BLOCK_ID) continue;

        uint32_t b_theta, b_E;
        decode_block(bid, b_theta, b_E);
        float E = get_rep_energy(b_E);

        if (E < b_E_trigger) {
            has_low_energy = true;
        }

        for (int i = 0; i < LOCAL_BINS; ++i) {
            total_weight += values_in[flat_index(cell, slot, i)];
        }
    }

    // □□ 4: □□ □□
    ActiveMask[cell] = (has_low_energy && total_weight > weight_active_min) ? 1 : 0;
}
```

1.3. K2: □□ □□ □□

1.3.1. □□

`src/cuda/kernels/k2_coarsetransport.cu`

1.3.2. □□□ □□

K2□ □□□□ □□□□ □□□□ □□□ □□□ □□□□□. □□ □□□□ □□ □□□ □□□□□□□ □□ □□□□ □□□□ 3-5□ □□ □□□ □□□□□.

1.3.3. □ □□□□

□□□□ □□(□□□ □□□□ □)□ □□ □□ □□□ □ □□□ □□□ □□□□. 150 MeV □□□ □□□□ 5% □□□ □□□□ □□□□ □□ □□□□□ □□□□ □□□□.
□□□□ □□□□ □□ □□□ □□□□ 3-5□ □□ □□□ □□ □ □□□□.

1.3.4. K3□□ □□ □□□

□□	K2 (□□)	K3 (□□)	□□
□□□ □□	□□□(□□□)	□(Vavilov)	3% □□□ □□
MCS □□□	□□□(□□□)	□(□□□ □□□)	5% □□ □□
□□ □□	□	□□	2-3□ □□ □□
□□□	5%	1 percent □□	□□□ □□ □□

Table 5: K2 vs K3 □□

1.3.5. □□□□ □□

□□ □□	**K2** □□□ □□
□□□ □□	□□: Vavilov □□□□ □□□ □□: □□□□ □□ • □□: 3% □□, 2□ □□
□□ □□ □□(MCS)	□□: □□□ □□ θ N(0, sigma²) □□□ □□: sigma² □□, □□ □ □□ □ □
□□ □□	□□: ds = min(□□□□, □□□□) □□: ds = 2-3□ □ □ • □□: □ □□ □□, □ □□ □□
□ □□	□□□ □□ (□□ □□) • w = **exp(-sigma_nuclear** ds)

Table 6: K2 □□□ □□□

1.3.6. □□□ □□

□□□□	□
□□□ □□	`(Nx * Nz + 255) / 256`
□□ □□	256 □□□
□□	□□ 1 □□□ (□□ □ □□□)
□□ □□	□□ □□ □□ 3-5□

Table 7: K2 □□□ □□

1.3.7. □□□□

```
__global__ void k2_coarsetransport(
    // □□ □□ □□ (□□ □□)
    const uint32_t* __restrict__ block_ids_in,
    const float* __restrict__ values_in,
    const uint8_t* __restrict__ ActiveMask,

    // □□□ □ □□
    const int Nx, const int Nz, const float dx, const float dz,
    const RLUT __restrict__ lut,

    // □□
    uint32_t* __restrict__ block_ids_out,
    float* __restrict__ values_out,
    double* __restrict__ EdepC,
    float* __restrict__ AbsorbedWeight_cutoff,
    float* __restrict__ AbsorbedWeight_nuclear,
    double* __restrict__ AbsorbedEnergy_nuclear,
    OutflowBucket* __restrict__ OutflowBuckets
);
```

1.3.8. □□□□ □□ □□

```
__device__ void coarse_transport_step(
    float& E, float& theta, float& x, float& z, float& w,
    float ds, const RLUT& lut
```

```
) {
    // □□□ □□ (□□□, □□ □□)
    float E_new = lut.lookup_E_inverse(lut.lookup_R(E) - ds);

    // □□ □□ □□ (□□ □□, □□□ □ □)
    float sigma_theta = highland_sigma(E, ds, X0_water);
    theta_variance += sigma_theta * sigma_theta;

    // □ □□ (□□□ □□)
    float sigma_nuc = Sigma_total(E);
    w *= exp(-sigma_nuc * ds);

    // □□ □□□□
    x += ds * sin(theta);
    z += ds * cos(theta);

    E = E_new;
}
```

1.4. K3: □□ □□ □□ (□□ □□)

1.4.1. □□

`src/cuda/kernels/k3_finetransport.cu`

1.4.2. □□□ □□

K3□ □□□□□□□ □□□□□□. □□□□ □□ □□□□ □□□□ □□□□□ □□□□ □□ □□□□ □□□□□□ □□□□ □□□□□□. □□□□ □□ □□□□ □□□□□ □□ □□ □□□□□ □□□□ □□ □ □□□□□□ □□□□□□□□□.

1.4.3. □ □□□□□

□□□□ □□□□ □□□□□ □□□□□ □□□□□ □□□□□. □□□□ □□ □□□□ □ □□ □□□□ □□□□□□. 10 MeV□□ 1mm □□ □□□□ □□□□ 20% □□□□ □ □□□□□. K3□ □□□□ □□□□ 1 percent □□ □□□□ □□□□□□.

1.4.4. □□ □□

□□	□□
□□ □ □□	• ActiveList□ □□ □□ • □ □□□□ □□□ □□ □ □□
RNG □□□	• □ □□□□ □□□ □□□ □□ □□□ • □□ □□□ □□□□ □□ □□
□□ □□	• □ □□□ □□ □□ “□□”□ □□ • □ □□ □□□
□□ □ □□	• □ □□□□ 512□□ □□ □□ □□ • □□□ □ □□ □□□ □□□
□□ □□ □□□	• □ □□□□ □□□ □□□ □□ • □□(θ), □□□(E), □□(x, z)
□ □ □□□	• □ □ □□ □□ □□□ • □□ □□□ □□ □□
□□ □□ □□	1. □□ □□ □□ (□□ + □□) 2. □□□ □□□ □□□ □□ 3. MCS □□□ 4. □ □□ 5. □□ □□□□ 6. □□ □□ □□
□□□ □□	• □□□ □□ □□ • EdepC, AbsorbedWeight, AbsorbedEnergy

Table 8: K3 □□□ □□ □□

1.4.5. □ □ □□□□ □□□

□ □ □□□ □□	□ □ □□□□□
□□□ □□ □□ □□□ □□□ □□□ □□ • □□□ □□□□□ (□□ □□) • □□□□□ □□	□ □□□ □ □ □□□ □□□□ □□ • □□ □□ □□ • □□□ □□

Table 9: □ □ □□□ □□

1.4.6. [illegible]

[illegible]	[illegible]
[illegible]	(n_active + 255) / 256
[illegible]	256 [illegible]
[illegible]	[illegible] 1 [illegible]
RNG [illegible]	[illegible] 1[illegible] ([illegible])
[illegible] [illegible]	[illegible] [illegible] 4 KB

Table 10: K3 [illegible]

1.4.7. [illegible]

```
__global__ void k3_finetransport(
    // [illegible]: [illegible] [illegible]
    const int* __restrict__ ActiveList,
    const int n_active,

    // [illegible] [illegible]
    const uint32_t* __restrict__ block_ids_in,
    const float* __restrict__ values_in,

    // [illegible] [illegible]
    const int Nx, const int Nz, const float dx, const float dz,
    const RLUT __restrict__ lut,
    const curandStateMRG32k3a* __restrict__ rng_states,

    // [illegible]
    uint32_t* __restrict__ block_ids_out,
    float* __restrict__ values_out,
    double* __restrict__ EdepC,
    float* __restrict__ AbsorbedWeight_cutoff,
    float* __restrict__ AbsorbedWeight_nuclear,
    double* __restrict__ AbsorbedEnergy_nuclear,
    OutflowBucket* __restrict__ OutflowBuckets
);
```

1.4.8. [illegible] ([illegible])

```
__global__ void k3_finetransport(...) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx >= n_active) return;

    int cell = ActiveList[idx];
    curandStateMRG32k3a local_rng = rng_states[idx];

    float cell_Edep = 0.0;
    float cell_E_nuc = 0.0;
    float w_cutoff = 0.0;

    for (int slot = 0; slot < Kb; ++slot) {
        uint32_t bid = block_ids_in[cell * Kb + slot];
        if (bid == EMPTY_BLOCK_ID) continue;

        uint32_t b_theta, b_E;
        decode_block(bid, b_theta, b_E);
```

```
    for (int lidx = 0; lidx < LOCAL_BINS; ++lidx) {
        float w = values_in[flat_index(cell, slot, lidx)];
        if (w < weight_epsilon) continue;

        // □□ □□ □□□
        int theta_local, E_local, x_sub, z_sub;
        decode_local_idx_4d(lidx, theta_local, E_local, x_sub, z_sub);

        float theta = get_theta_from_bins(b_theta, theta_local);
        float E = get_energy_from_bins(b_E, E_local);
        float x = cell_x + get_x_offset_from_bin(x_sub, dx);
        float z = cell_z + get_z_offset_from_bin(z_sub, dz);

        // □□ □□ □□
        while (true) {
            float ds = compute_max_step_physics(E, lut);
            ds = fminf(ds, compute_boundary_step(x, z, dx, dz, theta));

            // □□□ □□□ □□□ □□
            float dE_straggle = sample_energy_loss_with_straggling(E, ds, &local_rng);
            float E_new = lut.lookup_E_inverse(lut.lookup_R(E) - ds) + dE_straggle;
            E_new = fmaxf(E_new, E_cutoff);

            float dE = E - E_new;
            cell_Edep += w * dE;

            // MCS □□□
            float sigma_theta = highland_sigma(E, ds, X0_water);
            float dtheta = sample_mcs_angle(sigma_theta, &local_rng);
            theta += dtheta;

            // □ □□
            float sigma_nuc = Sigma_total(E);
            float w_nuc = w * (1.0f - exp(-sigma_nuc * ds));
            w -= w_nuc;
            cell_E_nuc += w_nuc * E;

            // □□ □□□□
            x += ds * sin(theta);
            z += ds * cos(theta);
            E = E_new;

            // □□ □□
            if (left_cell) {
                emit_to_bucket(OutflowBuckets, cell, face, theta, E, x, z, w);
                break;
            }

            if (E < E_cutoff) {
                w_cutoff += w;
                break;
            }
        }
    }
}

// □□□ □□
atomicAdd(&EdepC[cell], cell_Edep);
```

```
    atomicAdd(&AbsorbedWeight_cutoff[cell], w_cutoff);
    atomicAdd(&AbsorbedWeight_nuclear[cell], w - w_cutoff);
    atomicAdd(&AbsorbedEnergy_nuclear[cell], cell_E_nuc);

    rng_states[idx] = local_rng;
}
```

1.4.9. □ □ □□□□ □□

```
__device__ void sample_intra_bin(
    float& theta, float& E,
    int theta_local, int E_local,
    curandStateMRG32k3a* rng
) {
    // □ □ □□ □□□□ □□□□
    float u_theta = curand_uniform(rng) - 0.5f;  // [-0.5, 0.5]
    float u_E = curand_uniform(rng) - 0.5f;

    // □□ □□ □□
    theta += u_theta * dtheta_bin;
    E *= pow(10.0f, u_E * dlogE_bin);  // □□ □□
}
```

1.5. K4: □□ □□ □□

1.5.1. □□

`src/cuda/kernels/k4_transfer.cu`

1.5.2. □□□ □□

K4□ □□(K2/K3) □□ □□ □□ □□□ □ □□ □□□□□. □ □□□ 4□□ "□□"(□ □□□ □□: ±x, ±z)□ □□ □□□ □□□ □□□□□. K4□ □ □□□ □□ □□□ □□ □□ □□□ □□□□.

1.5.3. □ □□□□

□□□ □□□□□! K2/K3 □□ □ □□□ □ □□□ □□□ □ □□□□. □ □□ □□□ □□□□ □ □□ □□□ □□□. □□ □□ □□□ □□ □□□ □□□□ □□□□□.

1.5.4. □□ □□

□□	□□
□□	□ □□□□ 4□□ □□□□ □□: • □□ 0: +z □□□□□ □□□□ □□ • □□ 1: -z □□□□□ □□□□ □□ • □□ 2: +x □□□□□ □□□□ □□ • □□ 3: -x □□□□□ □□□□ □□
□□ □□	□□ (ix, iz)□ □□ □□: • +z: cell + Nx (iz+1 < Nz□□) • -z: cell - Nx (iz-1 >= 0□□) • +x: cell + 1 (ix+1 < Nx□□) • -x: cell - 1 (ix-1 >= 0□□)
□□ □□	□ □□ □□ □□: 1. 4□ □□□□ □□ □□ 2. □ □□□□ □□ □□ 3. block_id□ □□ □□□□ □□ • □: □□□□□ □□ □□□□ □□ • □□□□: □ □□ □□ 4. □□□□ □□ □□ (□□□□ □□)
□□	• □□ □□□□ □□□□ □□ □□ □□ • □□□□ □□ □□ □ (E, θ, x, z) • □□ □□ □□ □□

Table 11: K4 □□ □□ □□

1.5.5. □□□□ □□□□ □□□□

□□□□ □□ □□	**atomicCAS□**
□□ □□: • □□□□ A □□: slot[5] = EMPTY • □□□□ B □□: slot[5] = EMPTY • □□□□ A □□: slot[5] = BLOCK_42 • □□□□ B □□: slot[5] = BLOCK_99 • ⟶ □□□□ A□ □□□□ □□	□□□□ □□: • □□□□ A: atomicCAS ⟶ □□ • □□□□ B: atomicCAS ⟶ □□ • □□□□ B: slot[6] □□ • ⟶ □□ □□□□ □□

Table 12: □□□□ □□ □□□□ □□□□

1.5.6. □□□ □□

□□□□	□
□□□ □□	(Nx * Nz + 255) / 256
□□ □□	256 □□□
□□	□□ □□ 1 □□□
□□□ □□	□ (□□ □□, □□□ □□)
□□ □□□	□□ □□□ 1 KB

Table 13: K4 □□□ □□

1.5.7. □□□□

```
__global__ void k4_transfer(
    const OutflowBucket* __restrict__ OutflowBuckets,
    const int Nx, const int Nz,
    uint32_t* __restrict__ block_ids_out,
    float* __restrict__ values_out
);
```

1.5.8. □□□□

```
__global__ void k4_transfer(...) {
    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    if (cell >= Nx * Nz) return;

    int ix = cell % Nx;
    int iz = cell / Nx;

    // 4□ □□□□ □□
    int neighbors[4] = {
        iz + 1 < Nz ? cell + Nx : -1,  // +z
        iz - 1 >= 0 ? cell - Nx : -1,  // -z
        ix + 1 < Nx ? cell + 1 : -1,   // +x
        ix - 1 >= 0 ? cell - 1 : -1    // -x
    };

    for (int face = 0; face < 4; ++face) {
        int src_cell = neighbors[face];
        if (src_cell < 0) continue;

        const OutflowBucket& bucket = OutflowBuckets[src_cell * 4 + face];

        for (int k = 0; k < Kb_out; ++k) {
            uint32_t bid = bucket.block_id[k];
            if (bid == EMPTY_BLOCK_ID) continue;

            int slot = find_or_allocate_slot(block_ids_out, cell, bid);
            if (slot < 0) continue;

            for (int lidx = 0; lidx < LOCAL_BINS; ++lidx) {
                float w = bucket.value[k][lidx];
                if (w > 0) {
                    atomicAdd(&values_out[flat_index(cell, slot, lidx)], w);
                }
            }
```

```
        }
    }
}
```

1.5.9. □□□ □□ □□

```
__device__ int find_or_allocate_slot(
    uint32_t* block_ids,
    int cell,
    uint32_t bid
) {
    // □ □□ □□: □□ □□
    for (int slot = 0; slot < Kb; ++slot) {
        if (block_ids[cell * Kb + slot] == bid) {
            return slot;
        }
    }

    // □ □□ □□: □ □□ □□
    for (int slot = 0; slot < Kb; ++slot) {
        uint32_t expected = EMPTY_BLOCK_ID;
        uint32_t* ptr = &block_ids[cell * Kb + slot];
        if (atomicCAS(ptr, expected, bid) == expected) {
            return slot;
        }
    }

    return -1;  // □□ □□□ □□ □□
}
```

1.6. K5: □□ □□ □□

1.6.1. □□

`src/cuda/kernels/k5_audit.cu`

1.6.2. □□□ □□

K5□ □□□□□□□ □□□□□□□. □□ □ □□□□ □□□□□ □□□□□ □□□□□ □□□□□. □ □□ □□ □□□□ □□□□□□: “□□□□ □ = □□ □ + □□□□ □.”

1.6.3. □ □□□□□

□□ □□□□ □□□□□□. □□□□□ □□□□□ □□□□ □□□□□□□ □□□□ □□ □□□□□. K5□ □□□□ □□□□□□:

- □□□□□ □□ (□□ □□ □□)
- □□□□□ □□ □□□□ (□□ □□)
- □□□□ □□□□ □□ □□□□□ □□ □□

1.6.4. □□ □□

□	□□ (W_in)	□□+□□ (W_out+W_cut+W_nuc)	□□
Cell 0	1.0	0.7 + 0.2 + 0.1 = 1.0	□□ (0.0% □□)
Cell 1	0.8	0.6 + 0.15 + 0.05 = 0.8	□□ (0.0% □□)
Cell 2	0.5	0.4 + 0.08 + 0.02 = 0.5	□□ (0.0% □□)

Table 14: K5 □□ □□ □□

1.6.5. [illegible]

[illegible]	[illegible]
[illegible] [illegible]	(Nx * Nz + 255) / 256
[illegible] [illegible]	256 [illegible]
[illegible]	[illegible] 1 [illegible]
[illegible] [illegible]	[illegible] [illegible] ([illegible] [illegible] [illegible])

Table 15: K5 [illegible] [illegible]

1.6.6. [illegible]

```
__global__ void k5_audit(
    const uint32_t* __restrict__ block_ids_in,
    const float* __restrict__ values_in,
    const uint32_t* __restrict__ block_ids_out,
    const float* __restrict__ values_out,
    const float* __restrict__ AbsorbedWeight_cutoff,
    const float* __restrict__ AbsorbedWeight_nuclear,
    const double* __restrict__ AbsorbedEnergy_nuclear,
    const int Nx, const int Nz,
    AuditReport* __restrict__ reports
);
```

1.6.7. [illegible]

```
__global__ void k5_audit(...) {
    int cell = blockIdx.x * blockDim.x + threadIdx.x;
    if (cell >= Nx * Nz) return;

    // [illegible] [illegible] [illegible]
    float W_in = 0.0f;
    for (int slot = 0; slot < Kb; ++slot) {
        if (block_ids_in[cell * Kb + slot] != EMPTY_BLOCK_ID) {
            for (int i = 0; i < LOCAL_BINS; ++i) {
                W_in += values_in[flat_index(cell, slot, i)];
            }
        }
    }

    // [illegible] [illegible] [illegible]
    float W_out = 0.0f;
    for (int slot = 0; slot < Kb; ++slot) {
        if (block_ids_out[cell * Kb + slot] != EMPTY_BLOCK_ID) {
            for (int i = 0; i < LOCAL_BINS; ++i) {
                W_out += values_out[flat_index(cell, slot, i)];
            }
        }
    }

    // [illegible] [illegible]
    float W_cut = AbsorbedWeight_cutoff[cell];
    float W_nuc = AbsorbedWeight_nuclear[cell];

    // [illegible] [illegible]
    float W_expected = W_out + W_cut + W_nuc;
    float W_diff = fabsf(W_in - W_expected);
```

```
    float W_rel = W_diff / fmaxf(W_in, 1e-20f);

    // □□□ □□
    reports[cell].W_in = W_in;
    reports[cell].W_out = W_out;
    reports[cell].W_error = W_rel;
    reports[cell].pass = (W_rel < 1e-6f);
}
```

1.7. K6: □□ □□

1.7.1. □□

`src/cuda/kernels/k6_swap.cu`

1.7.2. □□□ □□

K6□ □□□ □□ □□□ □□□□ □□ □□ □□□ □□□□□. □ □□□ □□□ □□ □□□ □□□ □□□.

1.7.3. □ □□□□

K2-K4 □ □□□ “out” □□□ □□□□. □□ □□□ □□ □□□□ “in” □□□□□ □□□. 2.2 GB □□□□ □□□□ □□ □□□ □□□(□□)□ □□□□□.

1.7.4. □□ □□

□□ □	□□ □
□□□ □□□□: • block_ids_in ⟶ □□ A (□□ □□) • block_ids_out ⟶ □□ B (□ □□) • values_in ⟶ □□ C • values_out ⟶ □□ D	□□□ □□□□: • block_ids_in ⟶ □□ B (□□ □□) • block_ids_out ⟶ □□ A (□□ □□) • values_in ⟶ □□ D • values_out ⟶ □□ C
□□□ □□ □□!	□□ □□ □□ □□

Table 16: K6 □□□ □□ □□

1.7.5. □□□ □□□ □□

□□ **1:** □□□ □□	□□ **2:** □□□ □□
• cudaMemcpy(new_in, old_out, 2.2 GB) • □□□ 100 ms □□ • □□□ □□□ □□	• swap(in_ptr, out_ptr) • □□□ 0.001 microsec □□ • □□□ □□□ 0 □□ • 100,000□ □ □□!

Table 17: K6 □□ □□

1.7.6. □□

```
// □□□ □ □□ (□□ □□ □□)
void k6_swap_buffers(
    uint32_t*& block_ids_in,
    uint32_t*& block_ids_out,
    float*& values_in,
    float*& values_out
) {
    std::swap(block_ids_in, block_ids_out);
    std::swap(values_in, values_out);
}
```

1.7.7. □□□ □□ □□?

□□□ □□□ CPU □□□□□ - GPU □□□□ □□□ □□□ □□□□. □□ □□□ 2.2 GB □□□ □□□ □□□□.

1.8. □□□ □□ □□

1.8.1. □□ □□ □□

□□	□□
□□ □□ (□□)	□□□□ □□□□ □□□: • □□□ 0 □□: block_ids_in[0], values_in[0:31] • □□□ 1 □□: block_ids_in[1], values_in[32:63] • □□□ 2 □□: block_ids_in[2], values_in[64:95] GPU□ □□□ □□□ □□□□□□ □□: “block_ids_in[0:255] □ values_in[0:8191] □□” ⟶ □□ □□□□ 20-30□ □□
□□ □□ (□□)	• □□□ 0 □□: block_ids_in[0] • □□□ 1 □□: block_ids_in[1000] • □□□ 2 □□: block_ids_in[2000] GPU□ 3□□ □□ □□□ □□□□ □□ ⟶ 20-30□ □□

Table 18: □□□ □□ □□ □□

1.8.2. □□ □□□ □□

□□	□□ □□□	□□
K1	256 B	□□□ □□□ □□ □□ □□
K3	4 KB	□□ □ □□
K4	1 KB	□□ □□ □□

Table 19: □□ □□□ □□

1.9. □□ □□□ □□

□□	□□	□□
□□ □ □□	K2, K3	□ □ □□□ (60-90% □□)
□□/□□ □□	K2, K3	□□□□ 3-5□ □□ □□
□□□ □□	K4	□□□ □□ □□ □□
□ □ □□□	K3	□□ □□
□□□ □□	K6	2.2 GB □□□ □□ □□
□□ □□	□□	□□ □□□ □□□

Table 20: □□ □□□

1.10. □□: □□ □□ □□

1.10.1. □□□ □ □□ □□

```
// □□□ □□
dim3 grid( (Nx * Nz + 255) / 256 );
dim3 block(256);

// K1: ActiveMask
k1_activemask<<<grid, block>>>(...);

// K3: □□ □□ (□□ □□ □ □□ □□□)
dim3 grid_fine( (n_active + 255) / 256 );
k3_finetransport<<<grid_fine, block>>>(...);

// □□□
cudaDeviceSynchronize();
```

1.10.2. □□ □□□□ □□ □□

□□	□□
GPU vs CPU	**CPU □□:** □ □□ □□□□ □□□□ □□□□ □□□□ □□ **GPU □□:** 10,000□□ □□ □□□□ □□ □□□□ □□□□ □□
CUDA □□	□□ □□□□□ □□□□ □□□□ □□□□ • □ □□□□(□□□□)□ □□□□ □□□□(□□ □□) • □□ □□(□□□□)□ □□
□□□□ □□	□□□□: □□ □□ □□ □□□□ □□: □□□□ □ □□ □□□□ □□ (32-1024□) □□□□: □□□□ □□ □□ □□
□□□□ □□	□□ □□□□: □□ □□□□□ □□□□ □□□□ (□□) □□ □□□□: □□ □□□□ □□ □□□□□ □□ □□ □□ □□□□ (□□)
K1: ActiveMask	□□ □□ □□□□□ □□□□ □□□□ □□□□ □□ □□ • □□ □□ > 1000: □□ □□ • □□ □□ <= 1000: □□□□ □□
K2 vs K3	**K2 (□□):** □□□□ □□ □□□□ □□ □□ □□ **K3 (□□):** □□ □□□□□ □□ □□ □□□□□□ □□□□ □□ □□□□□□ K3□ □□ □□□□ □□
K4: □□	□□ □□□□ □□ □□ • □□□□ □□□□(□) □□ • 4□ □□ □□ □□□□ □□ □□ • □□□□ □□□□ □□
K5: □□	□□ □□ □□ □□ + □□ = □□ □□ + □□ □□□□ □□ □□□□ □□□□ □□
K6: □□	□□ □□ □□□□ □□ □□: □□ A□ □□ □□□□ □□ B□ □□ (□ □□) □□ □□: □□ A□ B□ □□ □□ (1□)

Table 21: CUDA □□□□ □□ □□ □□

1.11. □□

SM_2D CUDA □□□□□□□ □□□□ □□ □□□□ □□□□ □□□□□□ □□□□□ □□ 6□□ □□ □□□□□ □□□□□□:

- **K1 (ActiveMask)**: □□□□ □□□□□□□ □□□□ □ □□
- **K2 (□□)**: □□□□□ □□□□ □□ □□
- **K3 (□□)**: □□□□ □□ □□□□ □□□□ □□□□□□
- **K4 (□□)**: □ □ □□ □□
- **K5 (□□)**: □□ □□ □□
- **K6 (□□)**: □□ □□ □□

□ □□□□ □□ □□□□ □□, □□ □□□□ □□, □□□□□ □□ □□□□ □□□□□ GPU □□ □□□□ □□□□□□ □□□□□.

—

SM_2D CUDA □□□□□ □□

□□ 2.0 - □□□□□ □□□